



Workshop

Universidad
del País Vasco

Euskal Herriko
Unibertsitatea

Taller

Coding Deep Neural Networks for PDEs

Session III: Introducing Least Squares in (V)PINNs

Professors: Carlos Uriarte*, Ángel Javier Omella*, David Pardo*

Organizers: Ignacio Muga[†], Paulina Sepúlveda[†]

*Research Group on Applied Mathematical Modelling, Statistics, and Optimization (MATHMODE) at UPV/EHU

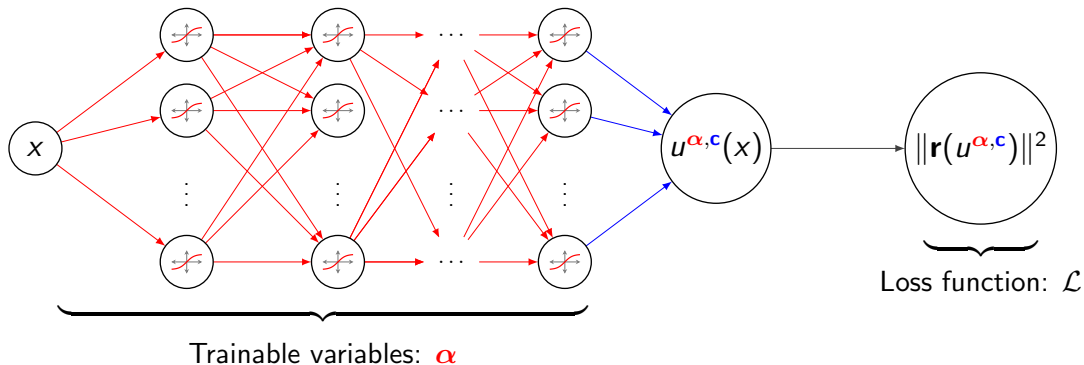
[†]IMA Numerics Research Group at PUCV

Follow us on: @mathmode.science*, @ima.numerics.pucv[†]

Biosfera Lodge, Olmué. January 19-24, 2025

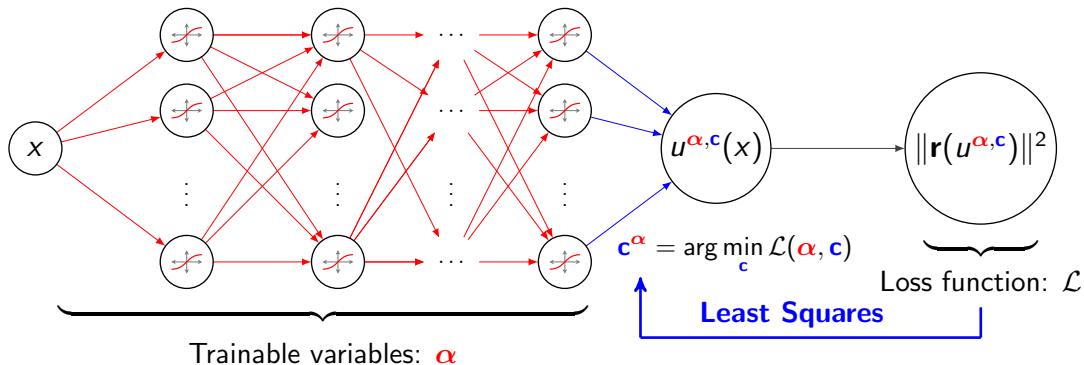
The LS/GD Optimizer

$$\underbrace{u^{\alpha, \mathbf{c}}(x)}_{\text{scalar-valued neural network}} = \mathbf{c} \cdot \mathbf{u}^{\alpha}(x) = \sum_{n=1}^N c_n \cdot u_n^{\alpha}(x), \quad \underbrace{\mathbf{u}^{\alpha}(x) := \{u_n^{\alpha}(x)\}_{n=1}^N}_{\text{vector-valued neural network}} \text{ spanning set to discretize } \mathbb{U}$$



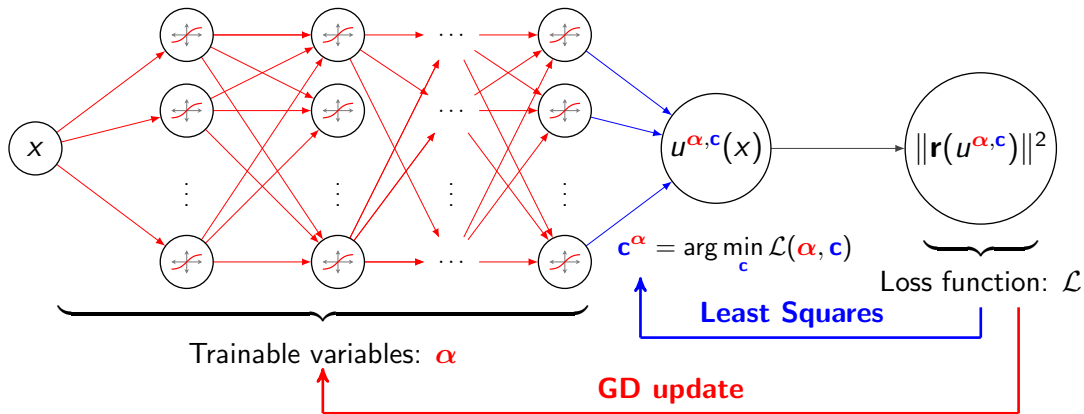
The LS/GD Optimizer

$$\underbrace{u^{\alpha, \mathbf{c}}(x)}_{\text{scalar-valued neural network}} = \mathbf{c} \cdot \mathbf{u}^{\alpha}(x) = \sum_{n=1}^N c_n \cdot u_n^{\alpha}(x), \quad \underbrace{\mathbf{u}^{\alpha}(x) := \{u_n^{\alpha}(x)\}_{n=1}^N}_{\text{vector-valued neural network}} \text{ spanning set to discretize } \mathbb{U}$$



The LS/GD Optimizer

$$\underbrace{u^{\alpha, \mathbf{c}}(x)}_{\text{scalar-valued neural network}} = \mathbf{c} \cdot \mathbf{u}^{\alpha}(x) = \sum_{n=1}^N c_n \cdot u_n^{\alpha}(x), \quad \underbrace{\mathbf{u}^{\alpha}(x) := \{u_n^{\alpha}(x)\}_{n=1}^N}_{\text{vector-valued neural network}} \text{ spanning set to discretize } \mathbb{U}$$



Petrov-Galerkin-Based Gradient-Descent Optimization

$$\mathcal{L}(\boldsymbol{\alpha}, \mathbf{c}) = \|\mathbf{r}(u^{\boldsymbol{\alpha}, \mathbf{c}})\|^2$$

$$\min_{\boldsymbol{\alpha}, \mathbf{c}} \mathcal{L}(\boldsymbol{\alpha}, \mathbf{c})$$

Conventional GD optimization

For $\boldsymbol{\theta} := (\boldsymbol{\alpha}, \mathbf{c})$

$$\boldsymbol{\theta}_t = \boldsymbol{\theta}_{t-1} - \eta_t \cdot \nabla_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}_{t-1})$$

η : learning rate

Petrov-Galerkin-Based Gradient-Descent Optimization

$$\mathcal{L}(\alpha, \mathbf{c}) = \|\mathbf{r}(u^{\alpha, \mathbf{c}})\|^2$$

$$\min_{\alpha, \mathbf{c}} \mathcal{L}(\alpha, \mathbf{c})$$

$$\min_{\alpha} \min_{\mathbf{c}} \mathcal{L}(\alpha, \mathbf{c})$$

Conventional GD optimization

For $\theta := (\alpha, \mathbf{c})$

$$\theta_t = \theta_{t-1} - \eta_t \cdot \nabla_{\theta} \mathcal{L}(\theta_{t-1})$$

η : learning rate

Hybrid LS/GD optimization

1. **Petrov-Galerkin.** Given α , find \mathbf{c}^{α} .
2. Update α using GD.

Implementation

Given α , we have:

$$\mathbf{r}(u^{\alpha, \mathbf{c}}) = \mathbf{B}^{\alpha} \mathbf{c} - \mathbf{I}$$

$$\mathbf{c}^{\alpha} = \arg \min_{\mathbf{c} \in \mathbb{R}^N} \|\mathbf{r}(u^{\alpha, \mathbf{c}})\|^2 = \arg \min_{\mathbf{c} \in \mathbb{R}^N} \|\mathbf{B}^{\alpha} \mathbf{c} - \mathbf{I}\|^2$$


Matrix and load vector (Poisson's equation: $-\Delta u = f$)

$$\mathbf{B}_{m,n}^{\alpha} = \int_{\alpha} \nabla u_n^{\alpha} \cdot \nabla v_m \approx \underbrace{\sum_{k=1}^K \gamma_k \nabla u_n^{\alpha}(x_k) \cdot \nabla v_m(x_k)}_{\text{Quadrature}}, \quad I_m = \int_{\alpha} f v_m \approx \underbrace{\sum_{k=1}^K \gamma_k f(x_k) v(x_k)}_{\text{Quadrature}}$$

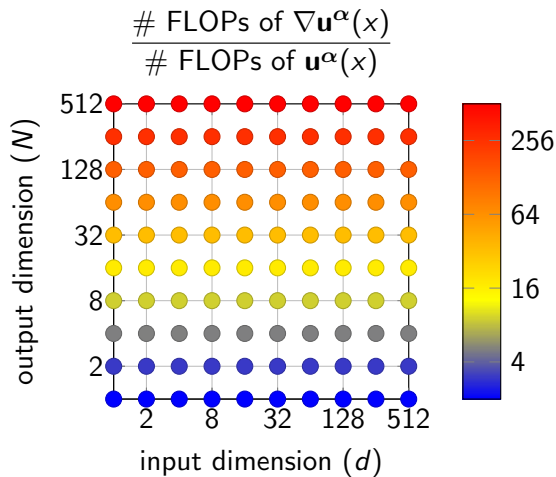
γ_k : quadrature weight, x_k : quadrature point, $1 \leq m \leq M$, $1 \leq n \leq N$.

- We compute \mathbf{c}^{α} via a direct LS solver, e.g. `tf.linalg.lstsq` ($\mathbf{B}^{\alpha}, \mathbf{I}$). **Cost:** $\mathcal{O}(MN^2)$

Automatic Differentiation



(a) Forward AD
(`tf.autodiff.ForwardAccumulator`).



(b) Backward AD
(`tf.GradientTape`).

Computational Cost

Using LS has the following dominant cost:

- Evaluation of $\nabla \mathbf{u}^\alpha$ over $\{x_k\}_{k=1}^K$ for $1 \leq n \leq N$.

$$\text{Cost: } \begin{cases} \mathcal{O}(KNC_{net}), & \text{if backward AD} \\ \mathcal{O}(KC_{net}), & \text{if forward AD} \end{cases}$$

- Summing for $1 \leq k, m, n \leq K, M, N$. Cost: $\mathcal{O}(KMN)$

- Solver of $\min_{\mathbf{c} \in \mathbb{R}^N} \|\mathbf{B}^\alpha \mathbf{c} - \mathbf{I}\|^2$. Cost: $\mathcal{O}(MN^2)$

Considerations

- $K \gtrsim M \gtrsim N$ (# of integration points \gtrsim highest test frequency \gtrsim # of spanning functions).
- $C_{net} \gtrsim MN$ (consider a network with “significant” evaluation cost).
- When applying **forward AD** (**backward AD**):

$$\text{Cost of LS: } \mathcal{O}(KNC_{net}) + \text{Cost of GD: } \mathcal{O}(KC_{net}) = \text{Cost of LS/GD: } \mathcal{O}(KNC_{net})$$

Computational Cost – Illustrative Experiment

- We consider a neural network with evaluation cost $C_{net} \gtrsim 512M$ and input dimension $d = 1$.
- We compare the costs of a **conventional GD optimizer** versus the **hybrid LS/GD optimizer** for $M = 5N$, $K = 10N$, and $N \in \{2, 4, 8, \dots, 512\}$.

$$\frac{\text{\# FLOPs when using the hybrid LS/GD optimizer}}{\text{\# FLOPs when using a conventional GD optimizer}}$$

N	2	4	8	16	32	64	128	256	512
Forward AD	1.33	1.33	1.33	1.33	1.34	1.34	1.33	1.35	1.37
Backward AD	1.49	1.83	2.49	3.83	6.46	11.82	22.35	43.77	86.05

Numerical Experiments

- Given f , find $u^* \in \mathbb{U} = H_0^1(\alpha)$ s.t.

$$\underbrace{\int_{\alpha} \nabla u^* \cdot \nabla v \, dx}_{b(u^*, v)} = \underbrace{\int_{\alpha} f v \, dx}_{l(v)}, \quad \forall v \in \mathbb{V} = H_0^1(\alpha)$$

- We equip $\mathbb{U} = H_0^1(\alpha) = \mathbb{V}$ with $b(\cdot, \cdot)$ as the inner product.
- (Continuum-level) residual $r(u)$ coincides with error function $e(u) = u - u^*$. Then,

$$\|e(u)\|_{\mathbb{U}} = \|r(u)\|_{\mathbb{V}}, \quad \forall u \in \mathbb{U} \quad (\mu = \alpha = 1).$$

- We consider a Fourier basis for the test space.

$$e(u) = r(u) = \sum_{m=1}^{\infty} \{b(u, v_m) - l(v_m)\} v_m \Rightarrow \|e(u)\|_{\mathbb{U}}^2 = \|r(u)\|_{\mathbb{V}}^2 = \sum_{m=1}^{\infty} \{b(u, v_m) - l(v_m)\}^2.$$

- VPINN loss: $\mathcal{L}(u) = \|r_M\|_{\mathbb{V}}^2 = \sum_{m=1}^M \{b(u, v_m) - l(v_m)\}^2.$

Experiment #1: Simple smooth solution in 1D

$$u^* = \sin(4x) \sin\left(\frac{x}{2}\right)$$

Discretization setup





- Fully-connected NN with 3 hidden layers:

$$x \in \mathbb{R}^{d=1} \rightarrow \mathbb{R}^{16} \rightarrow \mathbb{R}^{16} \rightarrow \mathbb{R}^{N=16} \rightarrow \mathbb{R} \ni u^{\alpha, c}(x)$$

- Activation function: \tanh
- Fourier modes (test space discretization): $M = 32$
- Iterations: 1,000
- Learning rate: 10^{-3}

We train two independent and equally initialized neural networks employing either Adam or LS/Adam optimizers.

Experiment #1: Simple smooth solution in 1D



- **Relative errors:**
2.72% (Adam),
0.19% (LS/Adam).
- Huge difference between LS/Adam and Adam alone from the very first training iteration
- Significant decrease difference of first Fourier modes when using LS/Adam compared to Adam alone.

Experiment #2: Smooth solution with high frequency in 1D

$$u^* = \sin(40x) \sin\left(\frac{x}{2}\right)$$

Discretization setup





- Fully-connected NN with 3 hidden layers:

$$x \in \mathbb{R}^{d=1} \rightarrow \mathbb{R}^{64} \rightarrow \mathbb{R}^{64} \rightarrow \mathbb{R}^{N=64} \rightarrow \mathbb{R} \ni u^{\alpha,c}(x)$$

- Activation function: \tanh
- Fourier modes (test space discretization): $M = 128$
- Iterations: 1,000
- Learning rate: 10^{-3}





We train two independent and equally initialized neural networks employing either Adam or LS/Adam optimizers.

Experiment #2: Smooth solution with high freq. in 1D (V1)



- **Relative errors:**
99.99% (Adam),
0.79% (LS/Adam).
- LS/Adam is able to decrease the loss, while Adam alone is not. This occurs due to a “slow” spectral bias of Adam.
- LS/Adam accelerates the overcoming of Adam’s spectral bias.

Experiment #2: Smooth solution with high freq. in 1D (V2)



- **Relative errors:**
3.28% (Adam),
0.04% (LS/Adam).
- In the long-term run, Adam alone is able to “discover” (and therefore diminish) the most significant nodes.
- LS/Adam accelerates the overcoming of this spectral bias notably.

Experiment #3: Smooth solution with high freq. in 2D

$$u^* = \sin(10x) \sin(2y) e^{\frac{x+y}{2}}$$

Discretization setup



- Fully-connected NN with 3 hidden layers:

$$(x, y) \in \mathbb{R}^{d=2} \rightarrow \mathbb{R}^{N=128} (3 \text{ times}) \rightarrow \mathbb{R} \ni u^{\alpha, c}(x, y)$$

- Activation function: tanh
- Fourier modes: $M = 64 \times 16$
- Iterations: 1,000
- Learning rate: 10^{-3}



We train two independent and equally initialized neural networks employing either Adam or LS/Adam optimizers.

Experiment #3: Smooth solution with high freq. in 2D







- **Relative errors:**
29.24% (Adam),
1.51% (LS/Adam).

- We observe a similar behavior as in the 1D case.






Experiment #3: Smooth solution with high freq. in 2D



- **Relative errors:**
29.24% (Adam),
1.51% (LS/Adam).
- We observe a similar behavior as in the 1D case.

References

-  Cyr, E. C., Gulian, M. A., Patel, R. G., Perego, M., & Trask, N. A. (2020). Robust training and initialization of deep neural networks: An adaptive basis viewpoint. In Mathematical and Scientific Machine Learning (pp. 512-536). PMLR.
-  Uriarte, C., Bastidas, M., Pardo, D., Taylor, J. M., & Rojas, S. (2024). Optimizing variational physics-informed neural networks using least squares. arXiv preprint arXiv:2407.20417.
-  Baharlouei, S., Taylor, J. M., Uriarte, C., & Pardo, D. (2024). A Least-Squares-Based Neural Network (LS-Net) for Solving Linear Parametric PDEs. arXiv preprint arXiv:2410.15089.